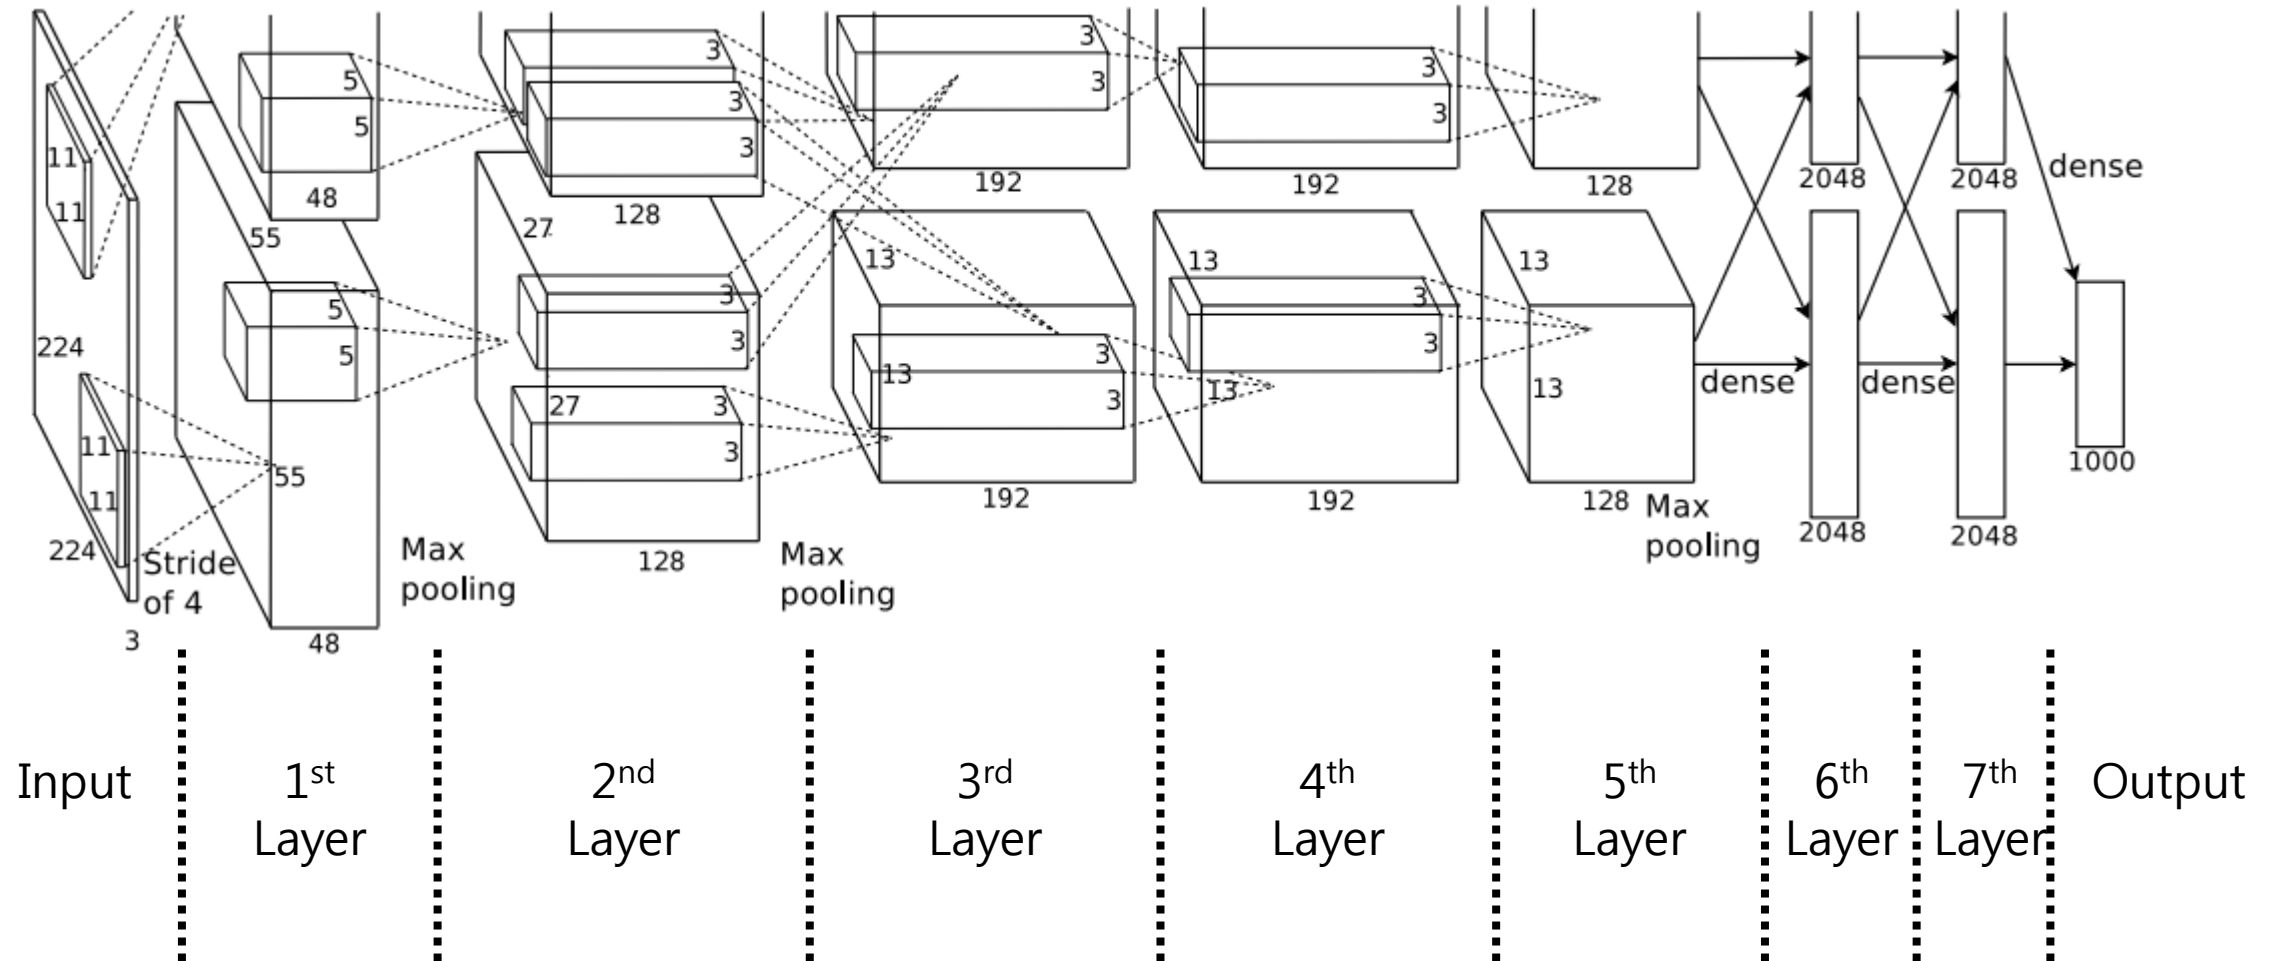


# AlexNet

Winner of ILSVRC 2012

# Overall Architecture



# Values from the paper

$$b_{x,y}^i = a_{x,y}^i / \left( k + \alpha \sum_{j=\max(0,i-n/2)}^{\min(N-1,i+n/2)} (a_{x,y}^j)^2 \right)^\beta$$

outputs computed using different kernels. The constants  $k$ ,  $n$ ,  $\alpha$ , and  $\beta$  are hyper-parameters whose values are determined using a validation set; we used  $k = 2$ ,  $n = 5$ ,  $\alpha = 10^{-4}$ , and  $\beta = 0.75$ . We applied this normalization after applying the ReLU nonlinearity in certain layers (see Section 3.5).

[17, 11, 4]). To be more precise, a pooling layer can be thought of as consisting of a grid of pooling units spaced  $s$  pixels apart, each summarizing a neighborhood of size  $z \times z$  centered at the location of the pooling unit. If we set  $s = z$ , we obtain traditional local pooling as commonly employed in CNNs. If we set  $s < z$ , we obtain overlapping pooling. This is what we use throughout our network, with  $s = 2$  and  $z = 3$ . This scheme reduces the top-1 and top-5 error rates by 0.4% and

# Values from the paper

connected layers are connected to all neurons in the previous layer. Response-normalization layers follow the first and second convolutional layers. Max-pooling layers, of the kind described in Section 3.4, follow both response-normalization layers as well as the fifth convolutional layer. The ReLU non-linearity is applied to the output of every convolutional and fully-connected layer.

factor of two during training. The recently-introduced technique, called “dropout” [10], consists of setting to zero the output of each hidden neuron with probability 0.5. The neurons which are

We use dropout in the first two fully-connected layers of Figure 2. Without dropout, our network exhibits substantial overfitting. Dropout roughly doubles the number of iterations required to converge.

# 1<sup>st</sup> Layer

The first convolutional layer filters the  $224 \times 224 \times 3$  input image with 96 kernels of size  $11 \times 11 \times 3$  with a stride of 4 pixels (this is the distance between the receptive field centers of neighboring

```
import conv2d from tensorflow.contrib.layers
import max_pool2d from tensorflow.contrib.layers
import local_response_normalization from tensorflow.nn
```

```
conv1 = conv2d(input,
               num_outputs=96,
               kernel_size=[11,11],
               stride=4,
               padding="VALID",
               activation_fn=tf.nn.relu)
```

```
lrn1 = local_response_normalization(conv1,
                                   bias=2
                                   alpha=0.0001,
                                   beta=0.75)
```

```
pool1 = max_pool2d(lrn1,
                   kernel_size=[3,3],
                   stride=2)
```

# 2<sup>nd</sup> Layer

neurons in a kernel map). The second convolutional layer takes as input the (response-normalized and pooled) output of the first convolutional layer and filters it with 256 kernels of size  $5 \times 5 \times 48$ .

```
conv2 = conv2d(lrn1,  
               num_outputs=256,  
               kernel_size=[5,5],  
               stride=1,  
               padding="VALID",  
               activation_fn=tf.nn.relu)
```

```
lrn2 = local_response_normalization(conv2,  
                                   bias=2  
                                   alpha=0.0001,  
                                   beta=0.75)
```

```
pool2 = max_pool2d(lrn2,  
                   kernel_size=[3,3],  
                   stride=2)
```

# 3<sup>rd</sup> Layer

pooling or normalization layers. The third convolutional layer has 384 kernels of size  $3 \times 3 \times 256$  connected to the (normalized, pooled) outputs of the second convolutional layer. The fourth

```
conv3 = conv2d(pool2,  
               num_outputs=384,  
               kernel_size=[3,3],  
               stride=1,  
               padding="VALID",  
               activation_fn=tf.nn.relu)
```

# 4<sup>th</sup> Layer

256 connected to the (normalized, pooled) outputs of the second convolutional layer. The fourth convolutional layer has 384 kernels of size  $3 \times 3 \times 192$  , and the fifth convolutional layer has 256 kernels of size  $3 \times 3 \times 192$ . The fully-connected layers have 4096 neurons each.

```
conv4 = conv2d(conv3,  
               num_outputs=384,  
               kernel_size=[3,3],  
               stride=1,  
               padding="VALID",  
               activation_fn=tf.nn.relu)
```



# 5<sup>th</sup> Layer

256 connected to the (normalized, pooled) outputs of the second convolutional layer. The fourth convolutional layer has 384 kernels of size  $3 \times 3 \times 192$  , and the fifth convolutional layer has 256 kernels of size  $3 \times 3 \times 192$ . The fully-connected layers have 4096 neurons each.

```
conv5 = conv2d(conv4,  
               num_outputs=256,  
               kernel_size=[3,3],  
               stride=1,  
               padding="VALID",  
               activation_fn=tf.nn.relu)
```

```
pool5 = max_pool2d(conv5,  
                   kernel_size=[3,3],  
                   stride=2)
```

# 6<sup>th</sup> ~ 7<sup>th</sup>, output Layers

256 connected to the (normalized, pooled) outputs of the second convolutional layer. The fourth convolutional layer has 384 kernels of size  $3 \times 3 \times 192$ , and the fifth convolutional layer has 256 kernels of size  $3 \times 3 \times 192$ . The fully-connected layers have 4096 neurons each.

```
import flatten from tensorflow.contrib.layers
import fully_connected from tensorflow.contrib.layers
```

```
flat = flatten(pool5)
```

```
fcl1 = fully_connected(flat,
                       num_outputs=4096,
                       activation_fn=tf.nn.relu)
```

```
fcl2 = fully_connected(fcl1,
                       num_outputs=4096,
                       activation_fn=tf.nn.relu)
```

```
out = fully_connected(fcl2,
                      num_outputs=1000,
                      activation_fn=None)
```